

Syed Faizan

## Magazine Subscription Analysis using two Group Classification Models

We implement logistic regression and Support Vector Machine Models to classify and predict the magazine subscription behaviour.

```
import pandas as pd
from sklearn.model_selection import train_test_split, GridSearchCV,
StratifiedKFold
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.metrics import accuracy_score, classification_report

# Set pandas options
pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', None)
```

```
file_path = r'C:\Users\sfaiz\OneDrive\Desktop\ALY 6020 Module 3
Project Syed Faizan\marketing_campaign.xlsx'
data = pd.read_excel(file_path)
```

```
data.head()
```

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome
Teenhome \						
0	5524	1957	Graduation	Single	58138.0	0
0						
1	2174	1954	Graduation	Single	46344.0	1
1						
2	4141	1965	Graduation	Together	71613.0	0
0						
3	6182	1984	Graduation	Together	26646.0	1
0						
4	5324	1981	PhD	Married	58293.0	1
0						

	Dt_Customer	Recency	MntWines	MntFruits	MntMeatProducts
MntFishProducts \					
0	2012-09-04	58	635	88	546
172					
1	2014-03-08	38	11	1	6
2					
2	2013-08-21	26	426	49	127
111					

```

3 2014-02-10      26      11      4      20
10
4 2014-01-19     94     173     43     118
46

```

```

      MntSweetProducts  MntGoldProds  NumDealsPurchases  NumWebPurchases
\
0      88      88      3      8
1      1      6      2      1
2     21     42      1      8
3      3      5      2      2
4     27     15      5      5

```

```

      NumCatalogPurchases  NumStorePurchases  NumWebVisitsMonth
AcceptedCmp3 \
0      10      4      7
0
1      1      2      5
0
2      2     10      4
0
3      0      4      6
0
4      3      6      5
0

```

```

      AcceptedCmp4  AcceptedCmp5  AcceptedCmp1  AcceptedCmp2  Complain \
0      0      0      0      0      0
1      0      0      0      0      0
2      0      0      0      0      0
3      0      0      0      0      0
4      0      0      0      0      0

```

```

      Z_CostContact  Z_Revenue  Response
0      3      11      1
1      3      11      0
2      3      11      0
3      3      11      0
4      3      11      0

```

```
data.shape
```

```
(2240, 29)
```

```
data.isnull().sum()
```

ID	0
Year_Birth	0
Education	0
Marital_Status	0
Income	24
Kidhome	0
Teenhome	0
Dt_Customer	0
Recency	0
MntWines	0
MntFruits	0
MntMeatProducts	0
MntFishProducts	0
MntSweetProducts	0
MntGoldProds	0
NumDealsPurchases	0
NumWebPurchases	0
NumCatalogPurchases	0
NumStorePurchases	0
NumWebVisitsMonth	0
AcceptedCmp3	0
AcceptedCmp4	0
AcceptedCmp5	0
AcceptedCmp1	0
AcceptedCmp2	0
Complain	0
Z_CostContact	0
Z_Revenue	0
Response	0

dtype: int64

```
data = data[data['Income'].notnull()]
```

```
data.shape
```

```
(2216, 29)
```

```
data.isnull().sum()
```

ID	0
Year_Birth	0
Education	0
Marital_Status	0
Income	0
Kidhome	0
Teenhome	0
Dt_Customer	0
Recency	0
MntWines	0
MntFruits	0

```
MntMeatProducts      0
MntFishProducts      0
MntSweetProducts     0
MntGoldProds         0
NumDealsPurchases    0
NumWebPurchases       0
NumCatalogPurchases  0
NumStorePurchases    0
NumWebVisitsMonth    0
AcceptedCmp3         0
AcceptedCmp4         0
AcceptedCmp5         0
AcceptedCmp1         0
AcceptedCmp2         0
Complain             0
Z_CostContact        0
Z_Revenue            0
Response             0
dtype: int64
```

```
data.dtypes
```

```
ID                int64
Year_Birth        int64
Education         object
Marital_Status    object
Income           float64
Kidhome          int64
Teenhome         int64
Dt_Customer       object
Recency          int64
MntWines         int64
MntFruits        int64
MntMeatProducts  int64
MntFishProducts  int64
MntSweetProducts int64
MntGoldProds     int64
NumDealsPurchases int64
NumWebPurchases  int64
NumCatalogPurchases int64
NumStorePurchases int64
NumWebVisitsMonth int64
AcceptedCmp3     int64
AcceptedCmp4     int64
AcceptedCmp5     int64
AcceptedCmp1     int64
AcceptedCmp2     int64
Complain         int64
Z_CostContact    int64
Z_Revenue        int64
```

```
Response          int64
dtype: object
```

**Ignore warnings as some Hyperparameter tuning may produce extensive warnings**

```
import warnings

# Ignore all warnings
warnings.filterwarnings("ignore")
```


## Examining the target variable

```
import matplotlib.pyplot as plt
import seaborn as sns

# Count the occurrences of each class in the target variable
response_counts = data['Response'].value_counts()

# Plot the imbalance
plt.figure(figsize=(8, 6))
sns.barplot(x=response_counts.index, y=response_counts.values,
            palette="viridis")

# Add labels and title
plt.title("Distribution of Target Variable 'Response'")
plt.xlabel("Response")
plt.ylabel("Count")
plt.xticks(ticks=[0, 1], labels=["No", "Yes"], rotation=0) #
# Customize if labels are binary
plt.show()
```



## Three issues to be dealt with by Data Cleansing :

1. The `Dt_Customer` column needs to be converted to an integer value that represents the time since subscription.
2. The 'object' columns need to be converted to numerical columns and encoded using ordinal encoding as one hot encoding may cause hyper-dimensionality.
3. The heavy imbalance in the target variable ought to be addressed and remedial measures ought to be taken.

`Dt_Customer` column converted to an integer value as 'Customer Tenure'.

```
# Dt_Customer column needs to be converted to an integer value  
data['Dt_Customer'] = pd.to_datetime(data['Dt_Customer'],
```

```
errors='coerce')
data = data.dropna(subset=['Dt_Customer'])
data['Customer_Tenure'] = (pd.Timestamp.now() -
data['Dt_Customer']).dt.days
```

```
data.head()
```

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome
0	5524	1957	Graduation	Single	58138.0	0
1	2174	1954	Graduation	Single	46344.0	1
2	4141	1965	Graduation	Together	71613.0	0
3	6182	1984	Graduation	Together	26646.0	1
4	5324	1981	PhD	Married	58293.0	1

	Dt_Customer	Recency	MntWines	MntFruits	MntMeatProducts
0	2012-09-04	58	635	88	546
1	2014-03-08	38	11	1	6
2	2013-08-21	26	426	49	127
3	2014-02-10	26	11	4	20
4	2014-01-19	94	173	43	118

	MntSweetProducts	MntGoldProds	NumDealsPurchases	NumWebPurchases
0	88	88	3	8
1	1	6	2	1
2	21	42	1	8
3	3	5	2	2
4	27	15	5	5

	NumCatalogPurchases	NumStorePurchases	NumWebVisitsMonth
0	10	4	7

1	1	2	5
0			
2	2	10	4
0			
3	0	4	6
0			
4	3	6	5
0			

	AcceptedCmp4	AcceptedCmp5	AcceptedCmp1	AcceptedCmp2	Complain	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	0	

	Z_CostContact	Z_Revenue	Response	Customer_Tenure
0	3	11	1	4461
1	3	11	0	3911
2	3	11	0	4110
3	3	11	0	3937
4	3	11	0	3959

Removing dt\_customer as it's data is found in customer tenure

```
# removing dt_customer as it's data is found in customer tenure
data_sansdt = data.drop(columns=['Dt_Customer'])
```

## Encode categorical variables

```
# Encode categorical variables
categorical_columns = ['Education', 'Marital_Status']
for col in categorical_columns:
    le = LabelEncoder()
    data_sansdt[col] = le.fit_transform(data_sansdt[col])

data_sansdt.head()
```

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome
Teenhome	\					
0	5524	1957	2	4	58138.0	0
0						
1	2174	1954	2	4	46344.0	1
1						



2	4141	1965	2	5	71613.0	0
0						
3	6182	1984	2	5	26646.0	1
0						
4	5324	1981	4	3	58293.0	1
0						

	Recency	MntWines	MntFruits	MntMeatProducts	MntFishProducts	\
0	58	635	88	546	172	
1	38	11	1	6	2	
2	26	426	49	127	111	
3	26	11	4	20	10	
4	94	173	43	118	46	

	MntSweetProducts	MntGoldProds	NumDealsPurchases	NumWebPurchases	\
0	88	88	3	8	
1	1	6	2	1	
2	21	42	1	8	
3	3	5	2	2	
4	27	15	5	5	

	NumCatalogPurchases	NumStorePurchases	NumWebVisitsMonth	\
AcceptedCmp3				
0	10	4	7	
0				
1	1	2	5	
0				
2	2	10	4	
0				
3	0	4	6	
0				
4	3	6	5	
0				

	AcceptedCmp4	AcceptedCmp5	AcceptedCmp1	AcceptedCmp2	Complain	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	0	

	Z_CostContact	Z_Revenue	Response	Customer_Tenure
0	3	11	1	4461
1	3	11	0	3911

2	3	11	0	4110
3	3	11	0	3937
4	3	11	0	3959

## Combine children columns into one

```
# Combine children columns into one
data_sansdt['Total_Children'] = data_sansdt['Kidhome'] +
data_sansdt['Teenhome']
data_sansdt.drop(['Kidhome', 'Teenhome', 'ID'], axis=1, inplace=True)
```

## Scale numerical features

```
# Scale numerical features
numerical_columns = [
    'Income', 'Recency', 'MntWines', 'MntFruits', 'MntMeatProducts',
    'MntFishProducts', 'MntSweetProducts', 'MntGoldProds',
    'NumDealsPurchases', 'NumWebPurchases', 'NumCatalogPurchases',
    'NumStorePurchases', 'NumWebVisitsMonth', 'Customer_Tenure'
]
scaler = StandardScaler()
data_sansdt[numerical_columns] =
scaler.fit_transform(data_sansdt[numerical_columns])

count = data[data['Year_Birth'] < 1940].shape[0]
print(count)

3
```

## Limiting dataset to those born only after 1940 by dropping three data points

```
data_sansdt = data_sansdt[data_sansdt['Year_Birth'] >= 1940]

from collections import Counter

Counter(data_sansdt['Response'])

Counter({0: 1880, 1: 333})

import seaborn as sns
```

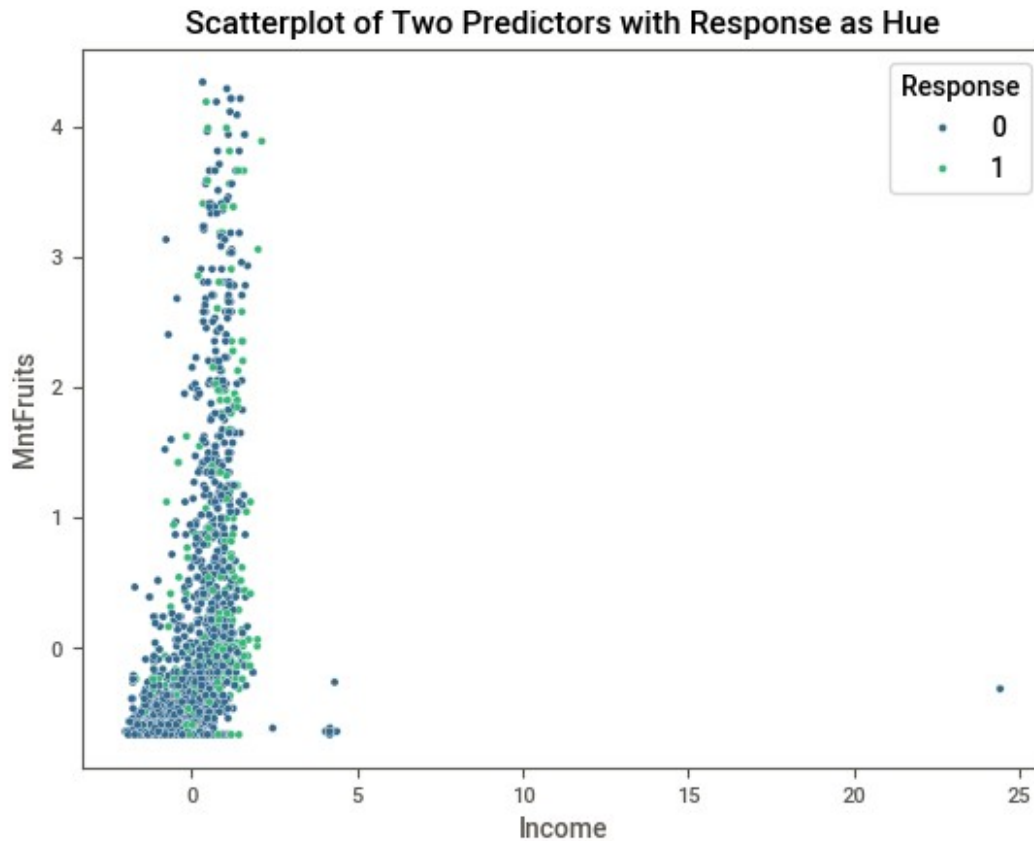
## Visually checking how target variable imbalance impacts the dataset by examining a scatterplot between two predictors

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Define the predictors (X) and target (y)
X = data_sansdt.drop(columns=['Response']) # Replace 'Response' with
the actual column name if different
y = data_sansdt['Response']

# Scatterplot for the first two predictors
sns.scatterplot(
    x=X.iloc[:, 3], # First predictor
    y=X.iloc[:, 6], # Second predictor
    hue=y,          # Target variable
    palette="viridis"
)

# Customize the plot
plt.title("Scatterplot of Two Predictors with Response as Hue")
plt.xlabel(X.columns[3]) # Label for the first predictor
plt.ylabel(X.columns[6]) # Label for the second predictor
plt.legend(title='Response')
plt.show()
```



```
data1 = data_sansdt
```

```
data1.head()
```

	Year_Birth	Education	Marital_Status	Income	Recency	MntWines
0	1957	2	4	0.234063	0.310532	0.978226
1	1954	2	4	-0.234559	-0.380509	-0.872024
2	1965	2	5	0.769478	-0.795134	0.358511
3	1984	2	5	-1.017239	-0.795134	-0.872024
4	1981	4	3	0.240221	1.554407	-0.391671

	MntFruits	MntMeatProducts	MntFishProducts	MntSweetProducts
0	1.549429	1.690227	2.454568	1.484827
1	-0.637328	-0.717986	-0.651038	-0.633880
2	0.569159	-0.178368	1.340203	-0.146821
3	-0.561922	-0.655551	-0.504892	-0.585174
4	0.418348	-0.218505	0.152766	-0.000703

MntGoldProds	NumDealsPurchases	NumWebPurchases		
NumCatalogPurchases \				
0	0.850031	0.351713	1.428553	
2.504712				
1	-0.732867	-0.168231	-1.125881	-
0.571082				
2	-0.037937	-0.688176	1.428553	-
0.229327				
3	-0.752171	-0.168231	-0.760962	-
0.912837				
4	-0.559135	1.391603	0.333796	
0.112428				

NumStorePurchases	NumWebVisitsMonth	AcceptedCmp3	AcceptedCmp4	\
0	-0.554143	0.693232	0	0
1	-1.169518	-0.131574	0	0
2	1.291982	-0.543978	0	0
3	-0.554143	0.280829	0	0
4	0.061232	-0.131574	0	0

AcceptedCmp5	AcceptedCmp1	AcceptedCmp2	Complain	
Z_CostContact \				
0	0	0	0	3
1	0	0	0	3
2	0	0	0	3
3	0	0	0	3
4	0	0	0	3

Z_Revenue	Response	Customer_Tenure	Total_Children	
0	11	1	1.529129	0
1	11	0	-1.188411	2
2	11	0	-0.205155	0
3	11	0	-1.059945	1
4	11	0	-0.951244	1

Transforming year birth into the numerical 'Age' column for better manipulation

```
data_sansdt['Age'] = (pd.Timestamp.now() -
pd.to_datetime(data_sansdt['Year_Birth'], format='%Y')).dt.days // 365
data_sansdt = data_sansdt.drop(columns = 'Year_Birth')
```

```
data_sansdt.head()
```

	Education	Marital_Status	Income	Recency	MntWines	MntFruits
0	2	4	0.234063	0.310532	0.978226	1.549429
1	2	4	-0.234559	-0.380509	-0.872024	-0.637328
2	2	5	0.769478	-0.795134	0.358511	0.569159
3	2	5	-1.017239	-0.795134	-0.872024	-0.561922
4	4	3	0.240221	1.554407	-0.391671	0.418348

	MntMeatProducts	MntFishProducts	MntSweetProducts	MntGoldProds
0	1.690227	2.454568	1.484827	0.850031
1	-0.717986	-0.651038	-0.633880	-0.732867
2	-0.178368	1.340203	-0.146821	-0.037937
3	-0.655551	-0.504892	-0.585174	-0.752171
4	-0.218505	0.152766	-0.000703	-0.559135

	NumDealsPurchases	NumWebPurchases	NumCatalogPurchases
0	0.351713	1.428553	2.504712
1	-0.168231	-1.125881	-0.571082
2	-0.688176	1.428553	-0.229327
3	-0.168231	-0.760962	-0.912837
4	1.391603	0.333796	0.112428

	NumWebVisitsMonth	AcceptedCmp3	AcceptedCmp4	AcceptedCmp5
0	0.693232	0	0	0
1	-0.131574	0	0	0
2	-0.543978	0	0	0
3	0.280829	0	0	0
4	-0.131574	0	0	0

	AcceptedCmp2	Complain	Z_CostContact	Z_Revenue	Response
0	0	0	3	11	1

1	0	0	3	11	0
2	0	0	3	11	0
3	0	0	3	11	0
4	0	0	3	11	0

	Customer_Tenure	Total_Children	Age
0	1.529129	0	67
1	-1.188411	2	70
2	-0.205155	0	59
3	-1.059945	1	40
4	-0.951244	1	43

## Scaling the Age column

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
data_sansdt['Age'] = scaler.fit_transform(data_sansdt[['Age']])
```

```
data_sansdt.head()
```

	Education	Marital_Status	Income	Recency	MntWines	MntFruits
0	2	4	0.234063	0.310532	0.978226	1.549429
1	2	4	-0.234559	-0.380509	-0.872024	-0.637328
2	2	5	0.769478	-0.795134	0.358511	0.569159
3	2	5	-1.017239	-0.795134	-0.872024	-0.561922
4	4	3	0.240221	1.554407	-0.391671	0.418348

	MntMeatProducts	MntFishProducts	MntSweetProducts	MntGoldProds
0	1.690227	2.454568	1.484827	0.850031
1	-0.717986	-0.651038	-0.633880	-0.732867
2	-0.178368	1.340203	-0.146821	-0.037937
3	-0.655551	-0.504892	-0.585174	-0.752171
4	-0.218505	0.152766	-0.000703	-0.559135

	NumDealsPurchases	NumWebPurchases	NumCatalogPurchases
0	0.351713	1.428553	2.504712
1	-0.168231	-1.125881	-0.571082
2	-0.688176	1.428553	-0.229327

3	-0.168231	-0.760962	-0.912837	-	
0.554143					
4	1.391603	0.333796	0.112428		
0.061232					
	NumWebVisitsMonth	AcceptedCmp3	AcceptedCmp4	AcceptedCmp5	
AcceptedCmp1 \					
0	0.693232	0	0	0	
0					
1	-0.131574	0	0	0	
0					
2	-0.543978	0	0	0	
0					
3	0.280829	0	0	0	
0					
4	-0.131574	0	0	0	
0					
	AcceptedCmp2	Complain	Z_CostContact	Z_Revenue	Response \
0	0	0	3	11	1
1	0	0	3	11	0
2	0	0	3	11	0
3	0	0	3	11	0
4	0	0	3	11	0
	Customer_Tenure	Total_Children	Age		
0	1.529129	0	1.018785		
1	-1.188411	2	1.275248		
2	-0.205155	0	0.334882		
3	-1.059945	1	-1.289387		
4	-0.951244	1	-1.032923		

## Scaling the newly encoded columns

```
# Scale emncoded features
encoded_columns = [
    'Education', 'Marital_Status', 'Z_CostContact',
    'Z_Revenue', 'Total_Children'
]
scaler = StandardScaler()
data_sansdt[encoded_columns] =
scaler.fit_transform(data_sansdt[encoded_columns])

data_sansdt.head()
```

	Education	Marital_Status	Income	Recency	MntWines	MntFruits
\						
0	-0.352454	0.254202	0.234063	0.310532	0.978226	1.549429



1	-0.352454	0.254202	-0.234559	-0.380509	-0.872024	-0.637328
2	-0.352454	1.182503	0.769478	-0.795134	0.358511	0.569159
3	-0.352454	1.182503	-1.017239	-0.795134	-0.872024	-0.561922
4	1.430358	-0.674098	0.240221	1.554407	-0.391671	0.418348

	MntMeatProducts	MntFishProducts	MntSweetProducts	MntGoldProds	\
0	1.690227	2.454568	1.484827	0.850031	
1	-0.717986	-0.651038	-0.633880	-0.732867	
2	-0.178368	1.340203	-0.146821	-0.037937	
3	-0.655551	-0.504892	-0.585174	-0.752171	
4	-0.218505	0.152766	-0.000703	-0.559135	

	NumDealsPurchases	NumWebPurchases	NumCatalogPurchases	
NumStorePurchases \				
0	0.351713	1.428553	2.504712	-
0.554143				
1	-0.168231	-1.125881	-0.571082	-
1.169518				
2	-0.688176	1.428553	-0.229327	
1.291982				
3	-0.168231	-0.760962	-0.912837	-
0.554143				
4	1.391603	0.333796	0.112428	
0.061232				

	NumWebVisitsMonth	AcceptedCmp3	AcceptedCmp4	AcceptedCmp5
AcceptedCmp1 \				
0	0.693232	0	0	0
0				
1	-0.131574	0	0	0
0				
2	-0.543978	0	0	0
0				
3	0.280829	0	0	0
0				
4	-0.131574	0	0	0
0				

	AcceptedCmp2	Complain	Z_CostContact	Z_Revenue	Response	\
0	0	0	0.0	0.0	1	
1	0	0	0.0	0.0	0	
2	0	0	0.0	0.0	0	
3	0	0	0.0	0.0	0	
4	0	0	0.0	0.0	0	

	Customer_Tenure	Total_Children	Age
0	1.529129	-1.264914	1.018785
1	-1.188411	1.404857	1.275248
2	-0.205155	-1.264914	0.334882
3	-1.059945	0.069971	-1.289387
4	-0.951244	0.069971	-1.032923

```
dfSummary(data_sansdt)
```

```
<pandas.io.formats.style.Styler at 0x1cf209a55e0>
```

```
dfSummary(data_sansdt)
```


```
<pandas.io.formats.style.Styler at 0x1cf209a55e0>
```

## Checking for the assumptions of Regression - Linearity

```
import seaborn as sns
import matplotlib.pyplot as plt

# Create a pair plot
sns.pairplot(data_sansdt, diag_kind="kde", corner=True)

# Show the plot
plt.show()
```



Dropping the Z Cost contact and Z Revenue variables due to there zero standard deviation and minimal information

```
columns_to_drop = ['Z_CostContact', 'Z_Revenue']
existing_columns = [col for col in columns_to_drop if col in
data_sansdt.columns]

if existing_columns:
    data_sansdt.drop(columns=existing_columns, inplace=True)
    print(f"Dropped columns: {existing_columns}")
else:
    print("No matching columns to drop.")
```

No matching columns to drop.

## Examining the skewed features and log transformation application

```
import numpy as np

# List of skewed predictors
skewed_predictors = [
    'Income', 'MntWines', 'MntFruits', 'MntMeatProducts',
    'MntFishProducts', 'MntSweetProducts', 'MntGoldProds'
]

# Handle negative values and apply log transformation
for col in skewed_predictors:
    # Replace negative values with NaN
    data_sansdt[col] = data_sansdt[col].apply(lambda x: np.nan if x <
0 else x)
    # Apply log transformation
    data_sansdt[col] = np.log1p(data_sansdt[col])

for col in skewed_predictors:
    data_sansdt[col].fillna(data_sansdt[col].median(), inplace=True)

# Check if transformation was successful
print("Log transformation applied to skewed predictors:")
print(data_sansdt[skewed_predictors].head())
```

Log transformation applied to skewed predictors:

	Income	MntWines	MntFruits	MntMeatProducts	MntFishProducts	\
0	0.210312	0.682200	0.935870	0.989626	1.239697	
1	0.511969	0.620372	0.646287	0.683673	0.671193	
2	0.570684	0.306389	0.450540	0.683673	0.850238	
3	0.511969	0.620372	0.646287	0.683673	0.671193	
4	0.215290	0.620372	0.349493	0.683673	0.142164	

	MntSweetProducts	MntGoldProds
0	0.910203	0.615202
1	0.654775	0.625582
2	0.654775	0.625582
3	0.654775	0.625582
4	0.654775	0.625582

## Checking skew after transformation

```
from scipy.stats import skew

for col in skewed_predictors:
    print(f"Skewness of {col}: {skew(data_sansdt[col]):.2f}")

Skewness of Income: 1.57
Skewness of MntWines: 0.58
Skewness of MntFruits: 1.04
Skewness of MntMeatProducts: 0.58
Skewness of MntFishProducts: 0.70
Skewness of MntSweetProducts: 0.95
Skewness of MntGoldProds: 0.89
```

## Examining the second assumption of regression modelling - Limited Multicollinearity

```
from scipy.cluster.hierarchy import linkage, dendrogram
import numpy as np


# Compute linkage for clustering
linkage_matrix = linkage(correlation_matrix, method='average')


# Plot dendrogram
plt.figure(figsize=(12, 8))
dendrogram(linkage_matrix, labels=correlation_matrix.columns,
            leaf_rotation=90)
plt.title("Dendrogram of Correlation Matrix")
plt.show()

# Reorder correlation matrix based on clustering
ordered_indices = dendrogram(linkage_matrix, no_plot=True)['leaves']
reordered_corr = correlation_matrix.iloc[ordered_indices,
                                         ordered_indices]

# Plot the reordered heatmap
plt.figure(figsize=(20, 14))
sns.heatmap(reordered_corr, annot=True, fmt=".2f", cmap="coolwarm",
            cbar=True, square=True)
plt.title("Clustered Correlation Matrix")
plt.show()
```

Dendrogram of Correlation Matrix





```

from statsmodels.stats.outliers_influence import
variance_inflation_factor
X = data_sansdt.drop(columns=['Response'])

# Replace missing values
X.fillna(X.median(), inplace=True)

# Drop constant columns
constant_cols = [col for col in X.columns if X[col].std() == 0]
X.drop(columns=constant_cols, inplace=True)

# Calculate VIF for each feature
vif_data = pd.DataFrame()
vif_data['Feature'] = X.columns

```

```
vif_data['VIF'] = [variance_inflation_factor(X.values, i) for i in
range(X.shape[1])]

print(vif_data)
```

	Feature	VIF
0	Education	1.090374
1	Marital_Status	1.009655
2	Income	9.953536
3	Recency	1.010645
4	MntWines	8.823181
5	MntFruits	8.648300
6	MntMeatProducts	9.890597
7	MntFishProducts	9.174705
8	MntSweetProducts	8.442964
9	MntGoldProds	6.763200
10	NumDealsPurchases	1.656132
11	NumWebPurchases	1.692045
12	NumCatalogPurchases	2.194465
13	NumStorePurchases	2.041905
14	NumWebVisitsMonth	2.349884
15	AcceptedCmp3	1.179194
16	AcceptedCmp4	1.392726
17	AcceptedCmp5	1.717832
18	AcceptedCmp1	1.420706
19	AcceptedCmp2	1.171125
20	Complain	1.015093
21	Customer_Tenure	1.268243
22	Total_Children	1.790751
23	Age	1.131812

## Data Splicing

```
# Split data into train-test sets
X = data_sansdt.drop('Response', axis=1)
y = data_sansdt['Response']
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, stratify=y, random_state=42)

X_train.head()
```

	Education	Marital_Status	Income	Recency	MntWines
MntFruits \					
2211	-0.352454	1.182503	0.889335	-0.242301	0.491196
0.752421					
925	0.538952	1.182503	0.511969	1.692615	0.620372
0.040490					
160	0.538952	0.254202	0.511969	-1.693488	0.620372
0.646287					
2014	-0.352454	0.254202	0.492917	-0.622374	1.113723



1.148251

725 -0.352454 -0.674098 0.511969 1.727167 0.620372

0.646287

	MntMeatProducts	MntFishProducts	MntSweetProducts	MntGoldProds
\				
2211	1.408681	0.698818	0.654775	0.160760
925	0.201308	0.671193	0.654775	0.863408
160	0.683673	0.671193	0.654775	0.625582
2014	0.749022	0.076677	0.412579	0.625582
725	0.683673	0.671193	0.654775	0.625582

	NumDealsPurchases	NumWebPurchases	NumCatalogPurchases	\
2211	-0.688176	0.698715	-0.229327	
925	0.351713	0.698715	-0.571082	
160	-0.688176	-1.125881	-0.912837	
2014	-0.688176	-0.760962	2.162957	
725	-0.688176	-0.760962	-0.912837	

	NumStorePurchases	NumWebVisitsMonth	AcceptedCmp3	AcceptedCmp4
\				
2211	0.984294	-1.368784	0	0
925	0.061232	1.105635	0	0
160	-1.169518	0.693232	0	0
2014	2.215044	-0.131574	0	0
725	-0.861830	0.280829	0	0

	AcceptedCmp5	AcceptedCmp1	AcceptedCmp2	Complain
Customer_Tenure	\			
2211	0	0	0	0
1.035240				-
925	0	0	0	0
1.548893				
160	0	0	0	0
0.046835				
2014	0	0	0	0
1.667476				
725	0	0	0	0
1.642981				-

Total\_Children

Age

2211	-1.264914	-1.032923
925	-1.264914	-0.861947
160	2.739742	0.933297
2014	-1.264914	0.933297
725	0.069971	0.249394

X\_test.head()

	Education	Marital_Status	Income	Recency	MntWines
MntFruits \					
1359	0.538952	0.254202	0.511969	-0.518718	0.620372
0.646287					
1779	-2.135266	-0.674098	0.419643	-1.658936	0.620372
0.646287					
1839	1.430358	-1.602398	0.184051	0.206876	1.059719
0.646287					
1151	-0.352454	1.182503	0.489537	1.623511	1.013451
1.254066					
561	-0.352454	0.254202	0.775604	0.897917	0.620372
1.431080					

	MntMeatProducts	MntFishProducts	MntSweetProducts	MntGoldProds
\				
1359	0.683673	0.671193	0.654775	0.625582
1779	0.683673	1.041494	0.919956	0.625582
1839	0.705934	0.531107	0.562035	1.266595
1151	0.121641	0.680486	1.121449	0.625582
561	1.492304	0.316259	1.282339	0.625582

	NumDealsPurchases	NumWebPurchases	NumCatalogPurchases	\
1359	0.351713	-0.760962	-0.571082	
1779	-0.688176	-0.760962	0.112428	
1839	4.511271	1.428553	0.795937	
1151	0.351713	2.158392	0.454182	
561	-1.208121	0.698715	1.137692	

	NumStorePurchases	NumWebVisitsMonth	AcceptedCmp3	AcceptedCmp4
\				
1359	-0.554143	0.280829	0	0
1779	2.215044	-1.781187	0	0
1839	0.984294	0.280829	0	0
1151	0.368919	0.280829	0	0

561	0.368919	-1.368784	0	0
AcceptedCmp5 AcceptedCmp1 AcceptedCmp2 Complain				
Customer_Tenure \				
1359	0	0	0	-
0.422559				
1779	0	0	0	-
0.106336				
1839	0	0	0	
1.440191				
1151	0	0	0	
1.558774				
561	0	0	0	
0.120949				
Total_Children Age				
1359	1.404857	-0.178045		
1779	-1.264914	-1.289387		
1839	0.069971	1.275248		
1151	0.069971	0.420370		
561	-1.264914	-0.947435		

## Logistic Regression Model

```
import statsmodels.api as sm
import pandas as pd

# Add a constant to the predictors (for intercept)
X_train_sm = sm.add_constant(X_train)

# Fit the logistic regression model
logit_model = sm.Logit(y_train, X_train_sm)
logit_result = logit_model.fit()

# Print the summary
print(logit_result.summary())
```

```
Optimization terminated successfully.
      Current function value: 0.272434
      Iterations 8
```

### Logit Regression Results

```
=====
=====
Dep. Variable:                Response   No. Observations:
1549
Model:                        Logit      Df Residuals:
```

1524						
Method:		MLE	Df	Model:		
24						
Date:		Thu, 21 Nov 2024		Pseudo R-squ.:		
0.3566						
Time:		21:32:07		Log-Likelihood:		
-422.00						
converged:		True		LL-Null:		
-655.90						
Covariance Type:		nonrobust		LLR p-value:		
7.882e-84						
=====						
=====						
		coef	std err	z	P> z	
-----						
-----						
const		-4.6890	0.603	-7.777	0.000	-
5.871	-3.507					
Education		0.2652	0.097	2.736	0.006	
0.075	0.455					
Marital_Status		-0.0709	0.090	-0.788	0.431	-
0.247	0.105					
Income		0.3629	0.479	0.758	0.449	-
0.576	1.301					
Recency		-0.7930	0.098	-8.067	0.000	-
0.986	-0.600					
MntWines		0.6878	0.420	1.639	0.101	-
0.135	1.510					
MntFruits		0.2425	0.367	0.661	0.509	-
0.476	0.961					
MntMeatProducts		1.2905	0.414	3.120	0.002	
0.480	2.101					
MntFishProducts		0.3377	0.366	0.922	0.357	-
0.380	1.056					
MntSweetProducts		-0.0958	0.360	-0.266	0.790	-
0.801	0.609					
MntGoldProds		-0.2865	0.330	-0.868	0.385	-
0.934	0.361					
NumDealsPurchases		0.2498	0.116	2.159	0.031	
0.023	0.477					
NumWebPurchases		0.3603	0.104	3.458	0.001	
0.156	0.565					
NumCatalogPurchases		0.1623	0.137	1.188	0.235	-
0.105	0.430					
NumStorePurchases		-0.7931	0.140	-5.684	0.000	-
1.067	-0.520					
NumWebVisitsMonth		0.0852	0.132	0.643	0.520	-
0.174	0.345					

AcceptedCmp3	2.0571	0.263	7.816	0.000	
1.541	2.573				
AcceptedCmp4	0.8895	0.332	2.675	0.007	
0.238	1.541				
AcceptedCmp5	1.4826	0.347	4.278	0.000	
0.803	2.162				
AcceptedCmp1	1.1984	0.355	3.374	0.001	
0.502	1.895				
AcceptedCmp2	1.1976	0.637	1.881	0.060	-
0.050	2.445				
Complain	0.4513	0.884	0.511	0.610	-
1.281	2.183				
Customer_Tenure	0.8997	0.110	8.197	0.000	
0.685	1.115				
Total_Children	-0.3735	0.133	-2.804	0.005	-
0.635	-0.112				
Age	0.0105	0.094	0.111	0.912	-
0.175	0.196				

```
=====
=====
```

```
logit_result.params.sort_values(ascending = False)
```


AcceptedCmp3	2.057070
AcceptedCmp5	1.482601
MntMeatProducts	1.290520
AcceptedCmp1	1.198401
AcceptedCmp2	1.197637
Customer_Tenure	0.899668
AcceptedCmp4	0.889458
MntWines	0.687785
Complain	0.451253
Income	0.362861
NumWebPurchases	0.360328
MntFishProducts	0.337670
Education	0.265183
NumDealsPurchases	0.249827
MntFruits	0.242478
NumCatalogPurchases	0.162336
NumWebVisitsMonth	0.085186
Age	0.010462
Marital_Status	-0.070877
MntSweetProducts	-0.095783
MntGoldProds	-0.286521
Total_Children	-0.373522
Recency	-0.793045
NumStorePurchases	-0.793099
const	-4.688960

dtype: float64

```
import matplotlib.pyplot as plt

# Assuming logit_result.params exists, create sorted coefficients
sorted_coefficients = logit_result.params.sort_values(ascending=False)

# Plot the coefficients as a colorful bar plot
plt.figure(figsize=(10, 6))
sorted_coefficients[0:12].plot(kind='bar', color='skyblue',
                                edgecolor='black')
plt.title('Logistic Regression Coefficients', fontsize=16)
plt.xlabel('Features', fontsize=14)
plt.ylabel('Coefficient Value', fontsize=14)
plt.xticks(rotation=45, fontsize=10, ha='right')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report

from sklearn.linear_model import LogisticRegression
logistic=LogisticRegression()
logistic.fit(X_train,y_train)
y_pred=logistic.predict(X_test)
```

## Logistic Regression Metrics

```
score=accuracy_score(y_pred,y_test)
print(f" The Logisitc Regression Model Metrics: \n The accuracy is :
{score}")
print(f"The Classification Report is : \n
{classification_report(y_pred,y_test)}")
print(f"The Confusion Matrix is : \n
{confusion_matrix(y_pred,y_test)}")
```

```
The Logisitc Regression Model Metrics:
The accuracy is : 0.8855421686746988
The Classification Report is :
```

	precision	recall	f1-score	support
0	0.97	0.90	0.94	608
1	0.40	0.71	0.51	56
accuracy			0.89	664
macro avg	0.69	0.81	0.72	664
weighted avg	0.92	0.89	0.90	664

```
The Confusion Matrix is :
```

```
[[548  60]
 [ 16  40]]
```

```
from sklearn.metrics import roc_curve, roc_auc_score
import matplotlib.pyplot as plt

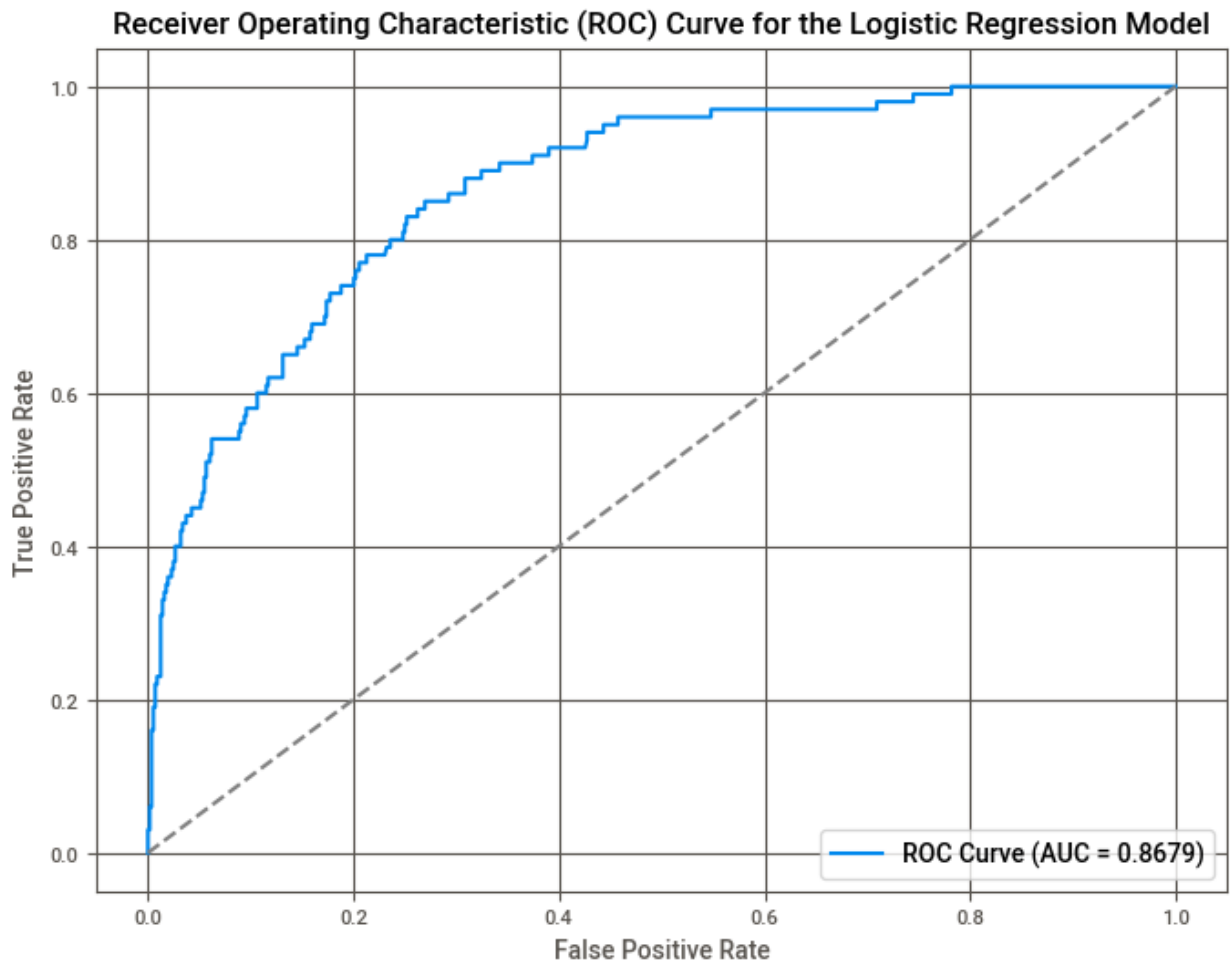
# Predict probabilities for the positive class
y_pred_proba = logistic.predict_proba(X_test)[: , 1]

# Calculate the ROC curve
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)

# Calculate the AUC
auc_score = roc_auc_score(y_test, y_pred_proba)
print(f"AUC Score: {auc_score:.4f}")

# Plot the ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f"ROC Curve (AUC = {auc_score:.4f})")
plt.plot([0, 1], [0, 1], linestyle='--', color='gray') # Diagonal line
plt.title("Receiver Operating Characteristic (ROC) Curve for the
Logistic Regression Model")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.grid()
plt.show()
```

AUC Score: 0.8679



Hyperparameter Tuning and rectifying the Target variable imbalance by assigning weights

```
## Hyperparameter tuning
from sklearn.linear_model import LogisticRegression
model=LogisticRegression()
penalty=['l1', 'l2', 'elasticnet']
c_values=[100,10,1.0,0.1,0.01]
solver=['newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga']
class_weight=[{0:w,1:y} for w in [1,10,50,100] for y in [1,10,50,100]]

params=dict(penalty=penalty,C=c_values,solver=solver,class_weight=clas
s_weight)

from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import StratifiedKFold
cv=StratifiedKFold()
```



```

grid=GridSearchCV(estimator=model,param_grid=params,scoring='accuracy',cv=cv)

grid.fit(X_train,y_train)

GridSearchCV(cv=StratifiedKFold(n_splits=5, random_state=None, shuffle=False),
             estimator=LogisticRegression(),
             param_grid={'C': [100, 10, 1.0, 0.1, 0.01],
                         'class_weight': [{0: 1, 1: 1}, {0: 1, 1: 10},
                                           {0: 1, 1: 50}, {0: 1, 1: 100},
                                           {0: 10, 1: 1}, {0: 10, 1: 10},
                                           {0: 10, 1: 50}, {0: 10, 1: 100},
                                           {0: 50, 1: 1}, {0: 50, 1: 10},
                                           {0: 50, 1: 50}, {0: 50, 1: 100},
                                           {0: 100, 1: 1}, {0: 100, 1: 10},
                                           {0: 100, 1: 50}, {0: 100, 1: 100}],
                         'penalty': ['l1', 'l2', 'elasticnet'],
                         'solver': ['newton-cg', 'lbfgs', 'liblinear', 'sag',
                                   'saga']},
             scoring='accuracy')

grid.best_params_

{'C': 1.0, 'class_weight': {0: 1, 1: 1}, 'penalty': 'l1', 'solver': 'saga'}


best_model = grid.best_estimator_

# Assuming coefficients are from a grid search fitted logistic regression model
grid_coefficients = pd.Series(best_model.coef_.flatten(),
                              index=X_train.columns).sort_values(ascending=False)

# Plot the top 12 coefficients as a colorful bar plot
plt.figure(figsize=(10, 6))
grid_coefficients[:12].plot(kind='bar', color='skyblue', edgecolor='black')
plt.title('Top 12 Logistic Regression Coefficients (GridSearchCV)',
          fontsize=16)
plt.xlabel('Features', fontsize=14)
plt.ylabel('Coefficient Value', fontsize=14)
plt.xticks(rotation=45, fontsize=10, ha='right')

```

```
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



```
y_pred1=grid.predict(X_test)
score=accuracy_score(y_pred1,y_test)
print(f" The Hyperparameter Tuned Logisitic Regression Model Metrics: \
n The accuracy is : {score}")
print(f"The Classification Report is : \n
{classification_report(y_pred1,y_test)}")
print(f"The Confusion Matrix is : \n
{confusion_matrix(y_pred1,y_test)}")
```

```
The Hyperparameter Tuned Logisitic Regression Model Metrics:
The accuracy is : 0.8900602409638554
The Classification Report is :
```

	precision	recall	f1-score	support
0	0.97	0.91	0.94	605
1	0.43	0.73	0.54	59
accuracy			0.89	664
macro avg	0.70	0.82	0.74	664
weighted avg	0.92	0.89	0.90	664

The Confusion Matrix is :

```
[[548  57]
 [ 16  43]]
```

```
from sklearn.metrics import roc_curve, roc_auc_score
import matplotlib.pyplot as plt
```

```
# Predict probabilities for the positive class
```

```
y_pred_proba1 = grid.predict_proba(X_test)[:, 1]
```

```
# Calculate the ROC curve
```

```
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba1)
```

```
# Calculate the AUC
```

```
auc_score = roc_auc_score(y_test, y_pred_proba1)
```

```
print(f"AUC Score: {auc_score:.4f}")
```

```
# Plot the ROC curve
```

```
plt.figure(figsize=(8, 6))
```

```
plt.plot(fpr, tpr, label=f"ROC Curve (AUC = {auc_score:.4f})")
```

```
plt.plot([0, 1], [0, 1], linestyle='--', color='gray') # Diagonal line
```

```
plt.title("Receiver Operating Characteristic (ROC) Curve of  
Hyperparameter tuned Logistic Model using GridsearchCV")
```

```
plt.xlabel("False Positive Rate")
```

```
plt.ylabel("True Positive Rate")
```

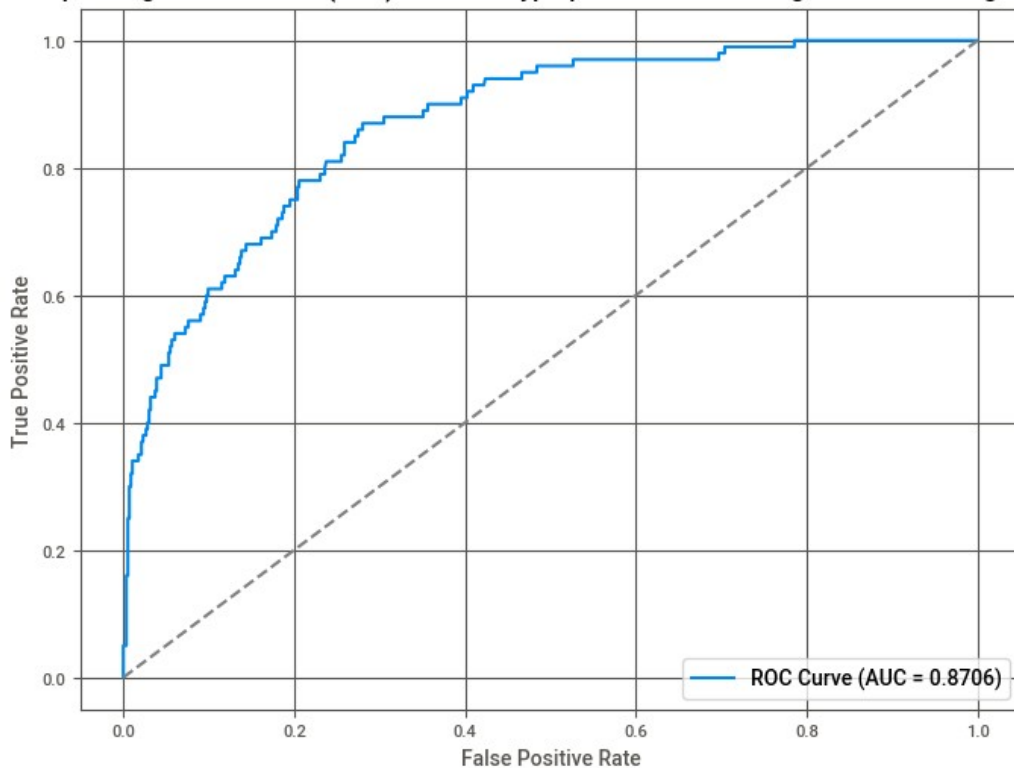
```
plt.legend(loc="lower right")
```

```
plt.grid()
```

```
plt.show()
```

AUC Score: 0.8706

Receiver Operating Characteristic (ROC) Curve of Hyperparameter tuned Logistic Model using GridsearchCV



```
from sklearn.metrics import roc_curve, roc_auc_score
import matplotlib.pyplot as plt
import numpy as np

# Predict probabilities for the positive class
y_pred_proba1 = grid.predict_proba(X_test)[: , 1]

# Calculate the ROC curve
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba1)

# Calculate the AUC
auc_score = roc_auc_score(y_test, y_pred_proba1)
print(f"AUC Score: {auc_score:.4f}")

# Plot the ROC curve with limited annotations
fig = plt.figure(figsize=(12, 8))
plt.plot(fpr, tpr, label=f"ROC Curve (AUC = {auc_score:.4f})",
marker='.')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray', label='Random
Model')

# Add limited annotations for thresholds (e.g., every 10th point)
for i, xy in enumerate(zip(fpr, tpr, thresholds)):
    if i % 10 == 0: # Annotate every 10th point
```

```
plt.annotate(f'{np.round(xy[2], 2)}', xy=(xy[0], xy[1]),
            fontsize=16, color='green')
```

```
# Axis labels and title
```

```
plt.title("Receiver Operating Characteristic (ROC) Curve of  
Hyperparameter tuned Logistic Regression Model with Threshold  
Annotations")
```

```
plt.xlabel("False Positive Rate")
```

```
plt.ylabel("True Positive Rate")
```

```
# Show legend
```

```
plt.legend(loc="lower right")
```

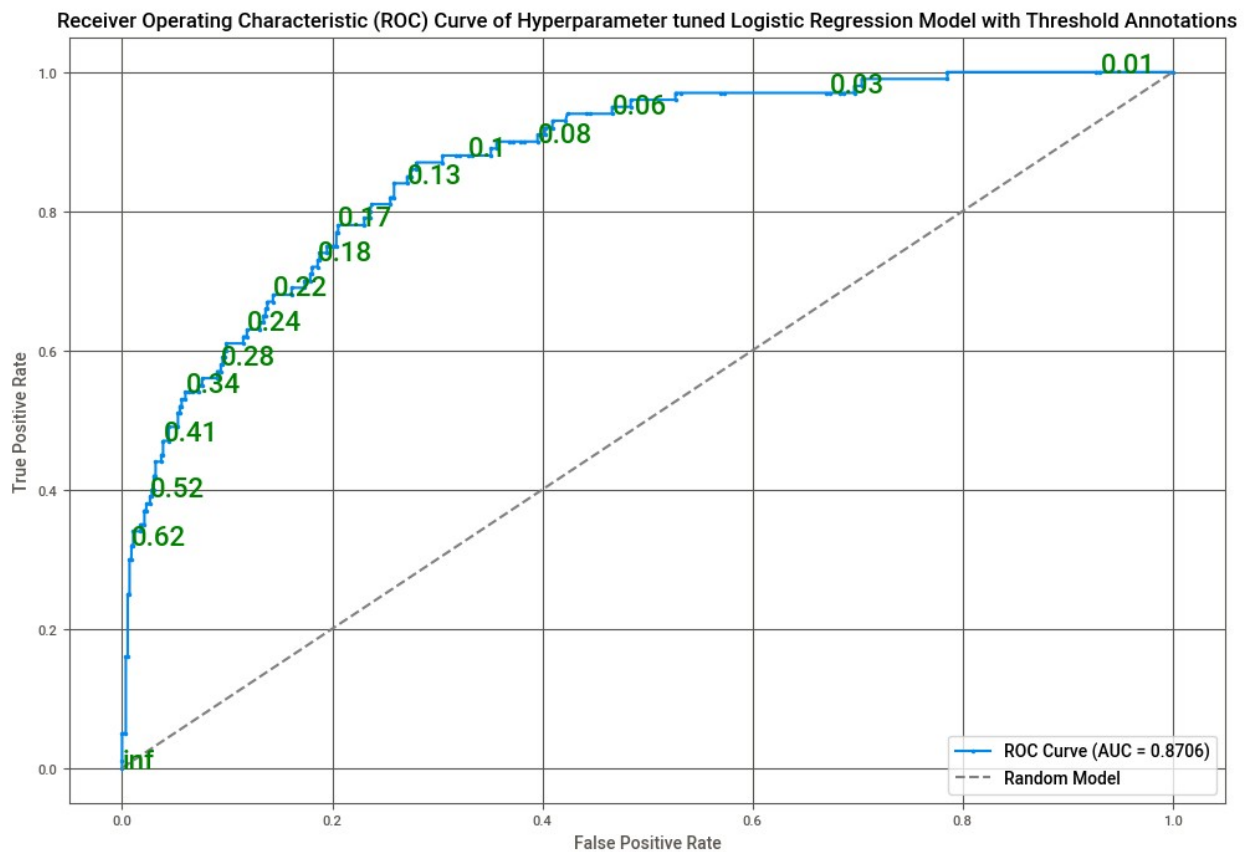
```
# Show grid
```

```
plt.grid()
```

```
# Show the plot
```

```
plt.show()
```

AUC Score: 0.8706



## Hyperparameter Tuning using randomized search CV and rectifying the Target variable imbalance by assigning weights

```
from sklearn.model_selection import RandomizedSearchCV

model=LogisticRegression()
randomcv=RandomizedSearchCV(estimator=model,param_distributions=params
,cv=5,scoring='accuracy')

randomcv.fit(X_train,y_train)

RandomizedSearchCV(cv=5, estimator=LogisticRegression(),
                    param_distributions={'C': [100, 10, 1.0, 0.1,
0.01],
                                        'class_weight': [{0: 1, 1: 1},
                                                         {0: 1, 1:
10}],
                                                         {0: 1, 1:
50}],
                                                         {0: 1, 1:
100}],
                                                         {0: 10, 1:
1}],
                                                         {0: 10, 1:
10}],
                                                         {0: 10, 1:
50}],
                                                         {0: 10, 1:
100}],
                                                         {0: 50, 1:
1}],
                                                         {0: 50, 1:
10}],
                                                         {0: 50, 1:
50}],
                                                         {0: 50, 1:
100}],
                                                         {0: 100, 1:
1}],
                                                         {0: 100, 1:
10}],
                                                         {0: 100, 1:
50}],
                                                         {0: 100, 1:
100}]],
                    'penalty': ['l1', 'l2',
'elasticnet'],
                    'solver': ['newton-cg',
```

```

'lbfgs',
                                                    'liblinear', 'sag',
                                                    'saga']},
scoring='accuracy')

randomcv.best_score_
0.8760622194383547

randomcv.best_params_
{'solver': 'liblinear',
 'penalty': 'l2',
 'class_weight': {0: 50, 1: 50},
 'C': 10}

y_pred2=randomcv.predict(X_test)

score=accuracy_score(y_pred2,y_test)
print(f" The Random Search CV Hyperparameter Tuned Logisitc Regression
Model Metrics: \n The accuracy is : {score}")
print(f"The Classification Report is : \n
{classification_report(y_pred2,y_test)}")
print(f"The Confusion Matrix is : \n
{confusion_matrix(y_pred2,y_test)}")

The Random Search CV Hyperparameter Tuned Logisitc Regression Model
Metrics:
The accuracy is : 0.8855421686746988
The Classification Report is :

```

	precision	recall	f1-score	support
0	0.97	0.90	0.94	606
1	0.41	0.71	0.52	58
accuracy			0.89	664
macro avg	0.69	0.80	0.73	664
weighted avg	0.92	0.89	0.90	664

```

The Confusion Matrix is :
[[547  59]
 [ 17  41]]

from sklearn.metrics import roc_curve, roc_auc_score
import matplotlib.pyplot as plt
import numpy as np

# Predict probabilities for the positive class
y_pred_proba2 = randomcv.predict_proba(X_test)[: , 1]

# Calculate the ROC curve
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba2)

```

```

# Calculate the AUC
auc_score = roc_auc_score(y_test, y_pred_proba2)
print(f"AUC Score: {auc_score:.4f}")

# Plot the ROC curve
plt.figure(figsize=(14, 8))
plt.plot(fpr, tpr, label=f"ROC Curve (AUC = {auc_score:.4f})")
plt.plot([0, 1], [0, 1], linestyle='--', color='gray', label="Random Model") # Diagonal line

# Add limited annotations for thresholds (e.g., every 50th point)
for i in range(0, len(thresholds), 10): # Change step size to adjust annotation density
    plt.annotate(f'{np.round(thresholds[i], 2)}', (fpr[i], tpr[i]),
        fontsize=14, color='green')

# Axis labels and title
plt.title("Receiver Operating Characteristic (ROC) of RandomCV Hyperparameter Tuned Logistic Regression Model with thresholds")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")

# Add legend
plt.legend(loc="lower right")

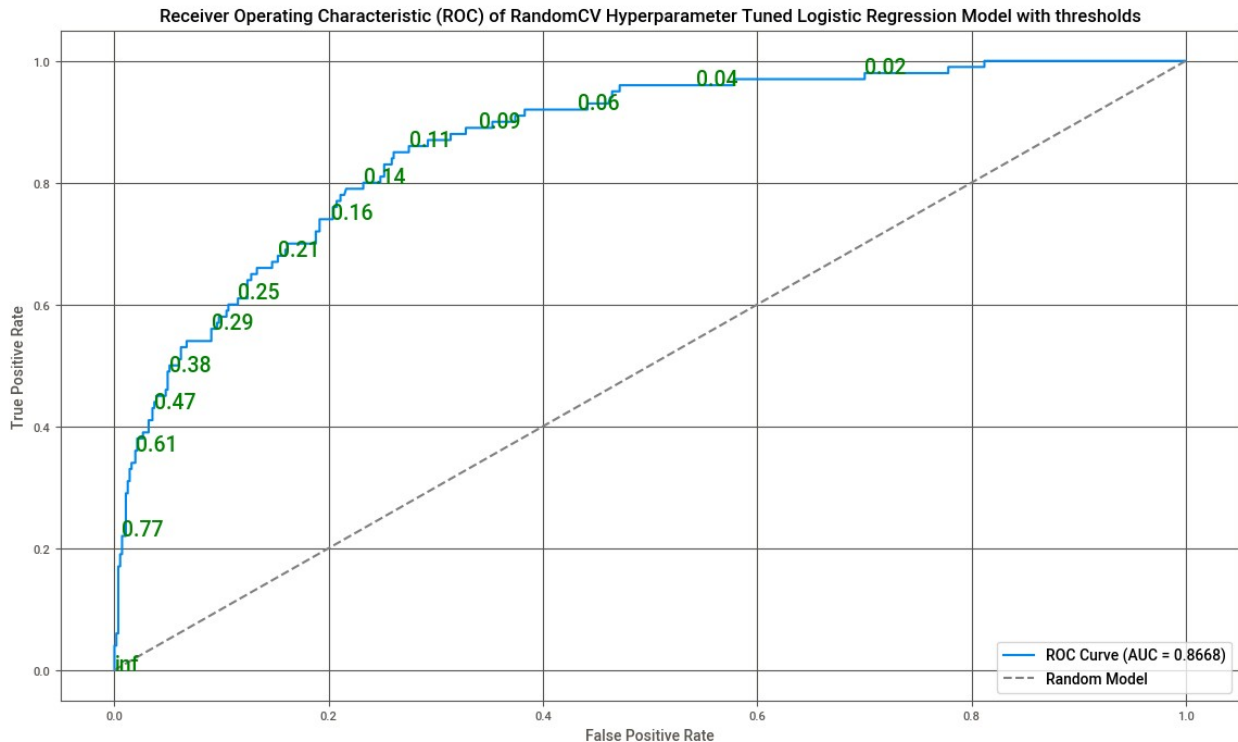
# Add grid
plt.grid()

# Show the plot
plt.show()

AUC Score: 0.8668

```





## Support Vector Machine

```
from sklearn.svm import SVC

svcl=SVC(kernel='linear', probability=True)

svcl.fit(X_train,y_train)

SVC(kernel='linear', probability=True)

# Get feature names and coefficients
coefficients = svcl.coef_[0] # Coefficients for the first class
features = X_train.columns   # Feature names (if X_train is a
                             # DataFrame)

# Create a DataFrame for better visualization
coef_df = pd.DataFrame({'Feature': features, 'Coefficient':
coefficients})
coef_df = coef_df.sort_values(by='Coefficient', ascending=False)

print("Coefficients for Predictors:")
print(coef_df)
```

Coefficients for Predictors:

	Feature	Coefficient
15	AcceptedCmp3	1.268511
17	AcceptedCmp5	0.961150
6	MntMeatProducts	0.928353

18	AcceptedCmp1	0.905359
4	MntWines	0.596882
16	AcceptedCmp4	0.513408
21	Customer_Tenure	0.441372
7	MntFishProducts	0.392129
2	Income	0.381166
19	AcceptedCmp2	0.268985
5	MntFruits	0.259621
20	Complain	0.211284
11	NumWebPurchases	0.192302
10	NumDealsPurchases	0.137807
0	Education	0.080210
12	NumCatalogPurchases	0.046844
23	Age	0.044956
14	NumWebVisitsMonth	-0.003640
9	MntGoldProds	-0.039200
1	Marital_Status	-0.049933
22	Total_Children	-0.188926
3	Recency	-0.355306
8	MntSweetProducts	-0.381656
13	NumStorePurchases	-0.417440

```
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report, roc_curve, roc_auc_score
import matplotlib.pyplot as plt
```

```
# Predictions
```

```
y_predl = svcl.predict(X_test)
```

```
# Accuracy
```

```
print("Accuracy of the SVM Linear Kernel is:", accuracy_score(y_test,
y_predl))
```

```
# Confusion Matrix
```

```
print("\nConfusion Matrix of the SVM Linear Kernel is:")
print(confusion_matrix(y_test, y_predl))
```

```
# Classification Report
```

```
print("\nClassification Report of the SVM Linear Kernel is:")
print(classification_report(y_test, y_predl))
```

```
# Predict probabilities for ROC curve
```

```
y_pred_proba = svcl.predict_proba(X_test)[:, 1]
```

```
# ROC Curve
```

```
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
auc_score = roc_auc_score(y_test, y_pred_proba)
```

```
# Print AUC
```

```
print(f"\nAUC Score: {auc_score:.4f}")
```

```

# Plot ROC Curve
plt.figure(figsize=(10, 6))
plt.plot(fpr, tpr, label=f"ROC Curve (AUC = {auc_score:.4f})",
marker='.')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray', label='Random
Model')
plt.title("Receiver Operating Characteristic (ROC) Curve for Linear
Kernel SVC")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.grid()
plt.show()

```

Accuracy of the SVM Linear Kernel is: 0.8689759036144579

Confusion Matrix of the SVM Linear Kernel is:

```

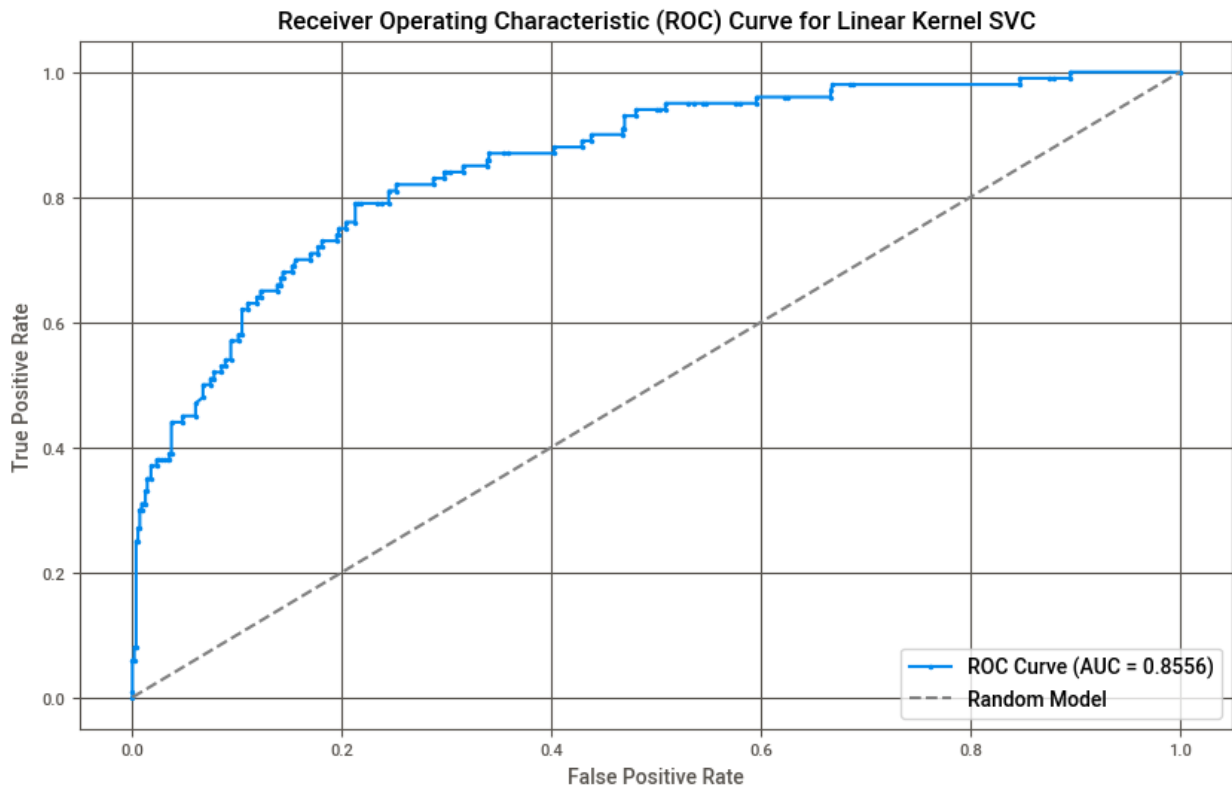
[[530  34]
 [ 53  47]]

```

Classification Report of the SVM Linear Kernel is:

	precision	recall	f1-score	support
0	0.91	0.94	0.92	564
1	0.58	0.47	0.52	100
accuracy			0.87	664
macro avg	0.74	0.70	0.72	664
weighted avg	0.86	0.87	0.86	664

AUC Score: 0.8556



## Radial Basis Function Kernel Support Vector Classifier

```
# Train SVC with RBF kernel and probability=True
rbf = SVC(kernel='rbf', probability=True)
rbf.fit(X_train, y_train)

SVC(probability=True)

# Predictions
y_pred_rbf = rbf.predict(X_test)

# Accuracy
print("Accuracy of the SVM RBF Kernel is:", accuracy_score(y_test,
y_pred_rbf))

# Confusion Matrix
print("\nConfusion Matrix of the SVM RBF Kernel is:")
print(confusion_matrix(y_test, y_pred_rbf))

# Classification Report
print("\nClassification Report of the SVM RBF Kernel is:")
print(classification_report(y_test, y_pred_rbf))

Accuracy of the SVM RBF Kernel is: 0.8795180722891566

Confusion Matrix of the SVM RBF Kernel is:
```

```
[[553  11]
 [ 69  31]]
```

Classification Report of the SVM RBF Kernel is:

	precision	recall	f1-score	support
0	0.89	0.98	0.93	564
1	0.74	0.31	0.44	100
accuracy			0.88	664
macro avg	0.81	0.65	0.68	664
weighted avg	0.87	0.88	0.86	664

```
# Predict probabilities for ROC curve
```

```
y_pred_proba_rbf = rbf.predict_proba(X_test)[:, 1]
```

```
# ROC Curve
```

```
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba_rbf)
```

```
auc_score = roc_auc_score(y_test, y_pred_proba_rbf)
```

```
# Print AUC
```

```
print(f"\nAUC Score: {auc_score:.4f}")
```

```
# Plot ROC Curve
```

```
plt.figure(figsize=(10, 6))
```

```
plt.plot(fpr, tpr, label=f"ROC Curve (AUC = {auc_score:.4f})",  
marker='.')
```

```
plt.plot([0, 1], [0, 1], linestyle='--', color='gray', label='Random  
Model')
```

```
plt.title("Receiver Operating Characteristic (ROC) Curve (RBF  
Kernel)")
```

```
plt.xlabel("False Positive Rate")
```


```
plt.ylabel("True Positive Rate")
```

```
plt.legend(loc="lower right")
```

```
plt.grid()
```

```
plt.show()
```

AUC Score: 0.8825



## Polynomial Kernel Support Vector Classifier

```
# Initialize SVC with Polynomial kernel and probability=True
poly = SVC(kernel='poly', probability=True)

# Fit the model to the training data
poly.fit(X_train, y_train)

SVC(kernel='poly', probability=True)

# Predict class labels on the test set
y_pred_poly = poly.predict(X_test)

# Predict probabilities for ROC curve
y_pred_proba_poly = poly.predict_proba(X_test)[:, 1]

# Calculate accuracy
accuracy_poly = accuracy_score(y_test, y_pred_poly)
print(f"Accuracy (Poly Kernel): {accuracy_poly:.4f}")

# Generate confusion matrix
conf_matrix_poly = confusion_matrix(y_test, y_pred_poly)
print("\nConfusion Matrix (Poly Kernel):")
print(conf_matrix_poly)

# Generate classification report
```

```
class_report_poly = classification_report(y_test, y_pred_poly)
print("\nClassification Report (Poly Kernel):")
print(class_report_poly)
```

Accuracy (Poly Kernel): 0.8795

Confusion Matrix (Poly Kernel):

```
[[549  15]
 [ 65  35]]
```

Classification Report (Poly Kernel):

	precision	recall	f1-score	support
0	0.89	0.97	0.93	564
1	0.70	0.35	0.47	100
accuracy			0.88	664
macro avg	0.80	0.66	0.70	664
weighted avg	0.86	0.88	0.86	664

*# Compute ROC curve*

```
fpr_poly, tpr_poly, thresholds_poly = roc_curve(y_test,
y_pred_proba_poly)
```

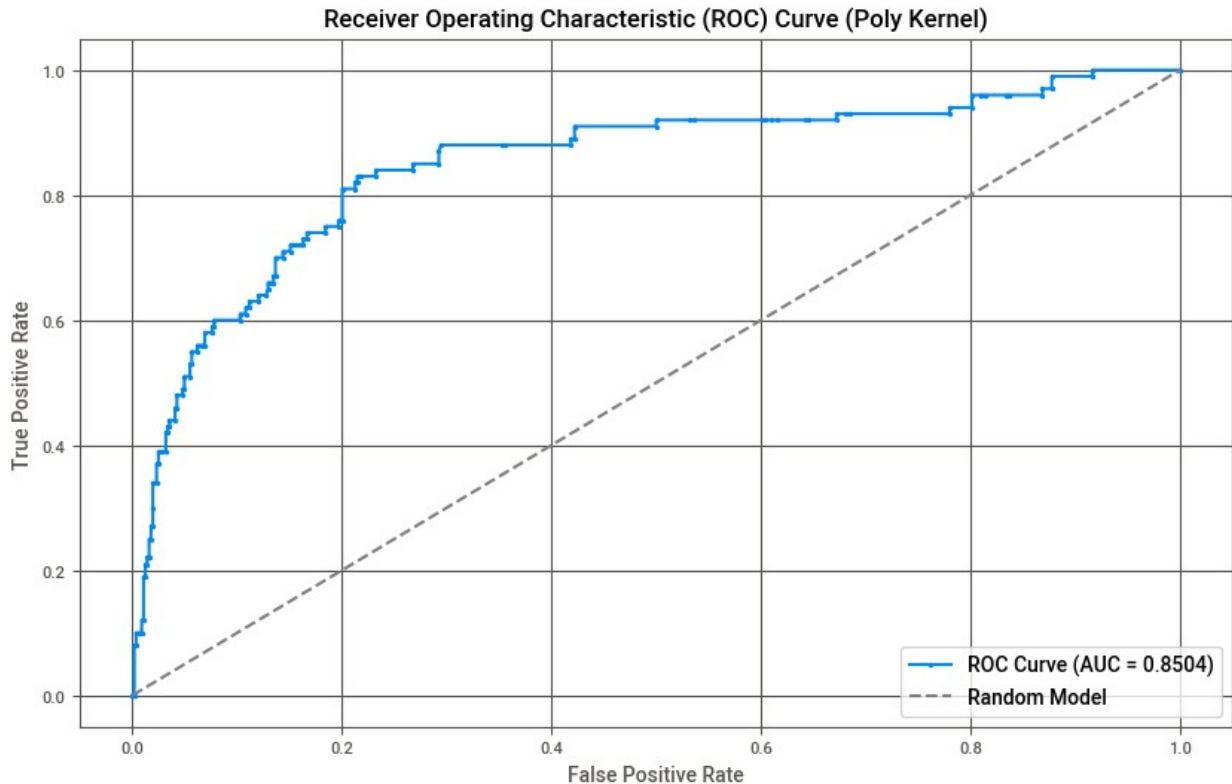
*# Compute AUC score*

```
auc_score_poly = roc_auc_score(y_test, y_pred_proba_poly)
print(f"\nAUC Score (Poly Kernel): {auc_score_poly:.4f}")
```

AUC Score (Poly Kernel): 0.8504

*# Plot ROC Curve*

```
plt.figure(figsize=(10, 6))
plt.plot(fpr_poly, tpr_poly, label=f"ROC Curve (AUC =
{auc_score_poly:.4f})", marker='.')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray', label='Random
Model')
plt.title("Receiver Operating Characteristic (ROC) Curve (Poly
Kernel)")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.grid()
plt.show()
```



## Sigmoid Kernel Support Vector Classifier

```
# Initialize SVC with Sigmoid kernel and probability=True
sigmoid = SVC(kernel='sigmoid', probability=True)

# Fit the model to the training data
sigmoid.fit(X_train, y_train)

SVC(kernel='sigmoid', probability=True)

# Predict class labels on the test set
y_pred_sigmoid = sigmoid.predict(X_test)

# Predict probabilities for ROC curve
y_pred_proba_sigmoid = sigmoid.predict_proba(X_test)[:, 1]

# Calculate accuracy
accuracy_sigmoid = accuracy_score(y_test, y_pred_sigmoid)
print(f"Accuracy (Sigmoid Kernel): {accuracy_sigmoid:.4f}")

# Generate confusion matrix
conf_matrix_sigmoid = confusion_matrix(y_test, y_pred_sigmoid)
print("\nConfusion Matrix (Sigmoid Kernel):")
print(conf_matrix_sigmoid)

# Generate classification report
```



```
class_report_sigmoid = classification_report(y_test, y_pred_sigmoid)
print("\nClassification Report (Sigmoid Kernel):")
print(class_report_sigmoid)
```

Accuracy (Sigmoid Kernel): 0.8223

Confusion Matrix (Sigmoid Kernel):

```
[[515  49]
 [ 69  31]]
```

Classification Report (Sigmoid Kernel):

	precision	recall	f1-score	support
0	0.88	0.91	0.90	564
1	0.39	0.31	0.34	100
accuracy			0.82	664
macro avg	0.63	0.61	0.62	664
weighted avg	0.81	0.82	0.81	664

*# Compute ROC curve*

```
fpr_sigmoid, tpr_sigmoid, thresholds_sigmoid = roc_curve(y_test,
y_pred_proba_sigmoid)
```

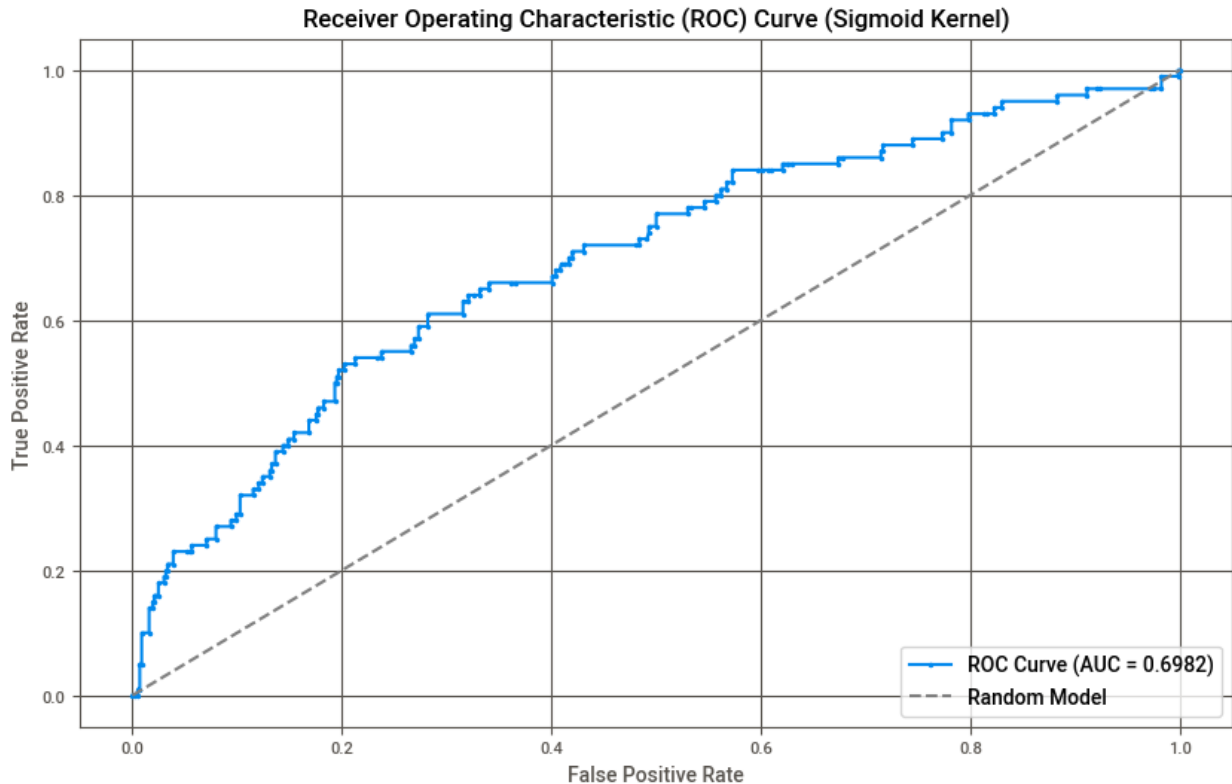
*# Compute AUC score*

```
auc_score_sigmoid = roc_auc_score(y_test, y_pred_proba_sigmoid)
print(f"\nAUC Score (Sigmoid Kernel): {auc_score_sigmoid:.4f}")
```

AUC Score (Sigmoid Kernel): 0.6982

*# Plot ROC Curve*

```
plt.figure(figsize=(10, 6))
plt.plot(fpr_sigmoid, tpr_sigmoid, label=f"ROC Curve (AUC =
{auc_score_sigmoid:.4f})", marker='.')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray', label='Random
Model')
plt.title("Receiver Operating Characteristic (ROC) Curve (Sigmoid
Kernel)")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.grid()
plt.show()
```



## Hyperparameter Tuning With SVC

```
from sklearn.model_selection import GridSearchCV

# defining parameter range
param_grid = {'C': [0.1, 1, 10, 100, 1000],
              'gamma': [1, 0.1, 0.01, 0.001, 0.0001],
              'kernel': ['rbf']}

gridsvc=GridSearchCV(SVC(probability=True),param_grid=param_grid,refit
=True,cv=5,verbose=3)

gridsvc.fit(X_train,y_train)
```

```
Fitting 5 folds for each of 25 candidates, totalling 125 fits
[CV 1/5] END .....C=0.1, gamma=1, kernel=rbf;; score=0.852 total
time= 0.9s
[CV 2/5] END .....C=0.1, gamma=1, kernel=rbf;; score=0.848 total
time= 0.9s
[CV 3/5] END .....C=0.1, gamma=1, kernel=rbf;; score=0.848 total
time= 0.9s
[CV 4/5] END .....C=0.1, gamma=1, kernel=rbf;; score=0.848 total
time= 0.8s
[CV 5/5] END .....C=0.1, gamma=1, kernel=rbf;; score=0.851 total
time= 0.8s
[CV 1/5] END .....C=0.1, gamma=0.1, kernel=rbf;; score=0.852 total
```

```
time= 0.3s
[CV 2/5] END .....C=0.1, gamma=0.1, kernel=rbf;, score=0.848 total
time= 0.3s
[CV 3/5] END .....C=0.1, gamma=0.1, kernel=rbf;, score=0.848 total
time= 0.3s
[CV 4/5] END .....C=0.1, gamma=0.1, kernel=rbf;, score=0.848 total
time= 0.3s
[CV 5/5] END .....C=0.1, gamma=0.1, kernel=rbf;, score=0.851 total
time= 0.3s
[CV 1/5] END .....C=0.1, gamma=0.01, kernel=rbf;, score=0.852 total
time= 0.2s
[CV 2/5] END .....C=0.1, gamma=0.01, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 3/5] END .....C=0.1, gamma=0.01, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 4/5] END .....C=0.1, gamma=0.01, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 5/5] END .....C=0.1, gamma=0.01, kernel=rbf;, score=0.851 total
time= 0.2s
[CV 1/5] END ....C=0.1, gamma=0.001, kernel=rbf;, score=0.852 total
time= 0.2s
[CV 2/5] END ....C=0.1, gamma=0.001, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 3/5] END ....C=0.1, gamma=0.001, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 4/5] END ....C=0.1, gamma=0.001, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 5/5] END ....C=0.1, gamma=0.001, kernel=rbf;, score=0.851 total
time= 0.2s
[CV 1/5] END ...C=0.1, gamma=0.0001, kernel=rbf;, score=0.852 total
time= 0.2s
[CV 2/5] END ...C=0.1, gamma=0.0001, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 3/5] END ...C=0.1, gamma=0.0001, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 4/5] END ...C=0.1, gamma=0.0001, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 5/5] END ...C=0.1, gamma=0.0001, kernel=rbf;, score=0.851 total
time= 0.2s
[CV 1/5] END .....C=1, gamma=1, kernel=rbf;, score=0.865 total
time= 0.8s
[CV 2/5] END .....C=1, gamma=1, kernel=rbf;, score=0.861 total
time= 0.8s
[CV 3/5] END .....C=1, gamma=1, kernel=rbf;, score=0.858 total
time= 0.9s
[CV 4/5] END .....C=1, gamma=1, kernel=rbf;, score=0.861 total
time= 0.9s
[CV 5/5] END .....C=1, gamma=1, kernel=rbf;, score=0.861 total
time= 0.8s
```

```

[CV 1/5] END .....C=1, gamma=0.1, kernel=rbf;, score=0.868 total
time= 0.3s
[CV 2/5] END .....C=1, gamma=0.1, kernel=rbf;, score=0.881 total
time= 0.3s
[CV 3/5] END .....C=1, gamma=0.1, kernel=rbf;, score=0.855 total
time= 0.3s
[CV 4/5] END .....C=1, gamma=0.1, kernel=rbf;, score=0.874 total
time= 0.3s
[CV 5/5] END .....C=1, gamma=0.1, kernel=rbf;, score=0.864 total
time= 0.3s
[CV 1/5] END .....C=1, gamma=0.01, kernel=rbf;, score=0.852 total
time= 0.2s
[CV 2/5] END .....C=1, gamma=0.01, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 3/5] END .....C=1, gamma=0.01, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 4/5] END .....C=1, gamma=0.01, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 5/5] END .....C=1, gamma=0.01, kernel=rbf;, score=0.851 total
time= 0.2s
[CV 1/5] END .....C=1, gamma=0.001, kernel=rbf;, score=0.852 total
time= 0.3s
[CV 2/5] END .....C=1, gamma=0.001, kernel=rbf;, score=0.848 total
time= 0.3s
[CV 3/5] END .....C=1, gamma=0.001, kernel=rbf;, score=0.848 total
time= 0.3s
[CV 4/5] END .....C=1, gamma=0.001, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 5/5] END .....C=1, gamma=0.001, kernel=rbf;, score=0.851 total
time= 0.2s
[CV 1/5] END .....C=1, gamma=0.0001, kernel=rbf;, score=0.852 total
time= 0.2s
[CV 2/5] END .....C=1, gamma=0.0001, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 3/5] END .....C=1, gamma=0.0001, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 4/5] END .....C=1, gamma=0.0001, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 5/5] END .....C=1, gamma=0.0001, kernel=rbf;, score=0.851 total
time= 0.2s
[CV 1/5] END .....C=10, gamma=1, kernel=rbf;, score=0.865 total
time= 0.9s
[CV 2/5] END .....C=10, gamma=1, kernel=rbf;, score=0.855 total
time= 0.9s
[CV 3/5] END .....C=10, gamma=1, kernel=rbf;, score=0.858 total
time= 0.9s
[CV 4/5] END .....C=10, gamma=1, kernel=rbf;, score=0.858 total
time= 0.9s
[CV 5/5] END .....C=10, gamma=1, kernel=rbf;, score=0.861 total

```

```

time= 0.8s
[CV 1/5] END .....C=10, gamma=0.1, kernel=rbf;, score=0.839 total
time= 0.4s
[CV 2/5] END .....C=10, gamma=0.1, kernel=rbf;, score=0.881 total
time= 0.4s
[CV 3/5] END .....C=10, gamma=0.1, kernel=rbf;, score=0.865 total
time= 0.4s
[CV 4/5] END .....C=10, gamma=0.1, kernel=rbf;, score=0.845 total
time= 0.4s
[CV 5/5] END .....C=10, gamma=0.1, kernel=rbf;, score=0.858 total
time= 0.4s
[CV 1/5] END .....C=10, gamma=0.01, kernel=rbf;, score=0.868 total
time= 0.2s
[CV 2/5] END .....C=10, gamma=0.01, kernel=rbf;, score=0.897 total
time= 0.3s
[CV 3/5] END .....C=10, gamma=0.01, kernel=rbf;, score=0.874 total
time= 0.2s
[CV 4/5] END .....C=10, gamma=0.01, kernel=rbf;, score=0.858 total
time= 0.2s
[CV 5/5] END .....C=10, gamma=0.01, kernel=rbf;, score=0.887 total
time= 0.2s
[CV 1/5] END .....C=10, gamma=0.001, kernel=rbf;, score=0.852 total
time= 0.2s
[CV 2/5] END .....C=10, gamma=0.001, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 3/5] END .....C=10, gamma=0.001, kernel=rbf;, score=0.848 total
time= 0.3s
[CV 4/5] END .....C=10, gamma=0.001, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 5/5] END .....C=10, gamma=0.001, kernel=rbf;, score=0.851 total
time= 0.2s
[CV 1/5] END ....C=10, gamma=0.0001, kernel=rbf;, score=0.852 total
time= 0.2s
[CV 2/5] END ....C=10, gamma=0.0001, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 3/5] END ....C=10, gamma=0.0001, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 4/5] END ....C=10, gamma=0.0001, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 5/5] END ....C=10, gamma=0.0001, kernel=rbf;, score=0.851 total
time= 0.2s
[CV 1/5] END .....C=100, gamma=1, kernel=rbf;, score=0.865 total
time= 0.8s
[CV 2/5] END .....C=100, gamma=1, kernel=rbf;, score=0.855 total
time= 0.8s
[CV 3/5] END .....C=100, gamma=1, kernel=rbf;, score=0.858 total
time= 0.8s
[CV 4/5] END .....C=100, gamma=1, kernel=rbf;, score=0.858 total
time= 0.9s

```

```
[CV 5/5] END .....C=100, gamma=1, kernel=rbf;, score=0.861 total
time= 0.9s
[CV 1/5] END .....C=100, gamma=0.1, kernel=rbf;, score=0.832 total
time= 0.4s
[CV 2/5] END .....C=100, gamma=0.1, kernel=rbf;, score=0.868 total
time= 0.4s
[CV 3/5] END .....C=100, gamma=0.1, kernel=rbf;, score=0.855 total
time= 0.4s
[CV 4/5] END .....C=100, gamma=0.1, kernel=rbf;, score=0.823 total
time= 0.4s
[CV 5/5] END .....C=100, gamma=0.1, kernel=rbf;, score=0.845 total
time= 0.4s
[CV 1/5] END .....C=100, gamma=0.01, kernel=rbf;, score=0.868 total
time= 0.4s
[CV 2/5] END .....C=100, gamma=0.01, kernel=rbf;, score=0.884 total
time= 0.4s
[CV 3/5] END .....C=100, gamma=0.01, kernel=rbf;, score=0.871 total
time= 0.4s
[CV 4/5] END .....C=100, gamma=0.01, kernel=rbf;, score=0.848 total
time= 0.3s
[CV 5/5] END .....C=100, gamma=0.01, kernel=rbf;, score=0.880 total
time= 0.4s
[CV 1/5] END ....C=100, gamma=0.001, kernel=rbf;, score=0.874 total
time= 0.3s
[CV 2/5] END ....C=100, gamma=0.001, kernel=rbf;, score=0.897 total
time= 0.3s
[CV 3/5] END ....C=100, gamma=0.001, kernel=rbf;, score=0.861 total
time= 0.3s
[CV 4/5] END ....C=100, gamma=0.001, kernel=rbf;, score=0.861 total
time= 0.3s
[CV 5/5] END ....C=100, gamma=0.001, kernel=rbf;, score=0.890 total
time= 0.3s
[CV 1/5] END ...C=100, gamma=0.0001, kernel=rbf;, score=0.852 total
time= 0.3s
[CV 2/5] END ...C=100, gamma=0.0001, kernel=rbf;, score=0.848 total
time= 0.3s
[CV 3/5] END ...C=100, gamma=0.0001, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 4/5] END ...C=100, gamma=0.0001, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 5/5] END ...C=100, gamma=0.0001, kernel=rbf;, score=0.848 total
time= 0.2s
[CV 1/5] END .....C=1000, gamma=1, kernel=rbf;, score=0.865 total
time= 0.8s
[CV 2/5] END .....C=1000, gamma=1, kernel=rbf;, score=0.855 total
time= 0.9s
[CV 3/5] END .....C=1000, gamma=1, kernel=rbf;, score=0.858 total
time= 0.8s
[CV 4/5] END .....C=1000, gamma=1, kernel=rbf;, score=0.858 total
```

```

time= 0.9s
[CV 5/5] END .....C=1000, gamma=1, kernel=rbf;, score=0.861 total
time= 0.9s
[CV 1/5] END .....C=1000, gamma=0.1, kernel=rbf;, score=0.835 total
time= 0.4s
[CV 2/5] END .....C=1000, gamma=0.1, kernel=rbf;, score=0.839 total
time= 0.4s
[CV 3/5] END .....C=1000, gamma=0.1, kernel=rbf;, score=0.842 total
time= 0.4s
[CV 4/5] END .....C=1000, gamma=0.1, kernel=rbf;, score=0.816 total
time= 0.4s
[CV 5/5] END .....C=1000, gamma=0.1, kernel=rbf;, score=0.854 total
time= 0.4s
[CV 1/5] END ....C=1000, gamma=0.01, kernel=rbf;, score=0.848 total
time= 1.2s
[CV 2/5] END ....C=1000, gamma=0.01, kernel=rbf;, score=0.887 total
time= 1.2s
[CV 3/5] END ....C=1000, gamma=0.01, kernel=rbf;, score=0.868 total
time= 1.3s
[CV 4/5] END ....C=1000, gamma=0.01, kernel=rbf;, score=0.842 total
time= 1.1s
[CV 5/5] END ....C=1000, gamma=0.01, kernel=rbf;, score=0.871 total
time= 1.3s
[CV 1/5] END ...C=1000, gamma=0.001, kernel=rbf;, score=0.861 total
time= 0.4s
[CV 2/5] END ...C=1000, gamma=0.001, kernel=rbf;, score=0.897 total
time= 0.4s
[CV 3/5] END ...C=1000, gamma=0.001, kernel=rbf;, score=0.865 total
time= 0.4s
[CV 4/5] END ...C=1000, gamma=0.001, kernel=rbf;, score=0.861 total
time= 0.4s
[CV 5/5] END ...C=1000, gamma=0.001, kernel=rbf;, score=0.896 total
time= 0.4s
[CV 1/5] END ..C=1000, gamma=0.0001, kernel=rbf;, score=0.874 total
time= 0.3s
[CV 2/5] END ..C=1000, gamma=0.0001, kernel=rbf;, score=0.900 total
time= 0.3s
[CV 3/5] END ..C=1000, gamma=0.0001, kernel=rbf;, score=0.858 total
time= 0.3s
[CV 4/5] END ..C=1000, gamma=0.0001, kernel=rbf;, score=0.868 total
time= 0.2s
[CV 5/5] END ..C=1000, gamma=0.0001, kernel=rbf;, score=0.893 total
time= 0.3s

```

```

GridSearchCV(cv=5, estimator=SVC(probability=True),
             param_grid={'C': [0.1, 1, 10, 100, 1000],
                          'gamma': [1, 0.1, 0.01, 0.001, 0.0001],
                          'kernel': ['rbf']},
             verbose=3)

```

```

# Print the best parameters
print("Best Parameters:", gridsvc.best_params_)

Best Parameters: {'C': 1000, 'gamma': 0.0001, 'kernel': 'rbf'}

# Access the best model
best_svc = gridsvc.best_estimator_
best_svc

SVC(C=1000, gamma=0.0001, probability=True)

## Prediction with hyperparameter tuned SVC model
y_pred4=gridsvc.predict(X_test)
# Accuracy
print("Accuracy of the hyperparameter tuned SVC is:",
accuracy_score(y_test, y_pred4))

# Confusion Matrix
print("\nConfusion Matrix of the hyperparameter tuned SVC is:")
print(confusion_matrix(y_test, y_pred4))

# Classification Report
print("\nClassification Report of the hyperparameter tuned SVC is:")
print(classification_report(y_test, y_pred4))

Accuracy of the hyperparameter tuned SVC is: 0.8870481927710844

Confusion Matrix of the hyperparameter tuned SVC is:
[[552  12]
 [ 63  37]]

Classification Report of the hyperparameter tuned SVC is:

```

	precision	recall	f1-score	support
0	0.90	0.98	0.94	564
1	0.76	0.37	0.50	100
accuracy			0.89	664
macro avg	0.83	0.67	0.72	664
weighted avg	0.88	0.89	0.87	664

```

# Use the best model to predict probabilities
y_pred_proba = best_svc.predict_proba(X_test)[:, 1]

from sklearn.metrics import roc_curve, roc_auc_score
import matplotlib.pyplot as plt

# Compute ROC curve and AUC score for the probabilities
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
auc_score = roc_auc_score(y_test, y_pred_proba)

```



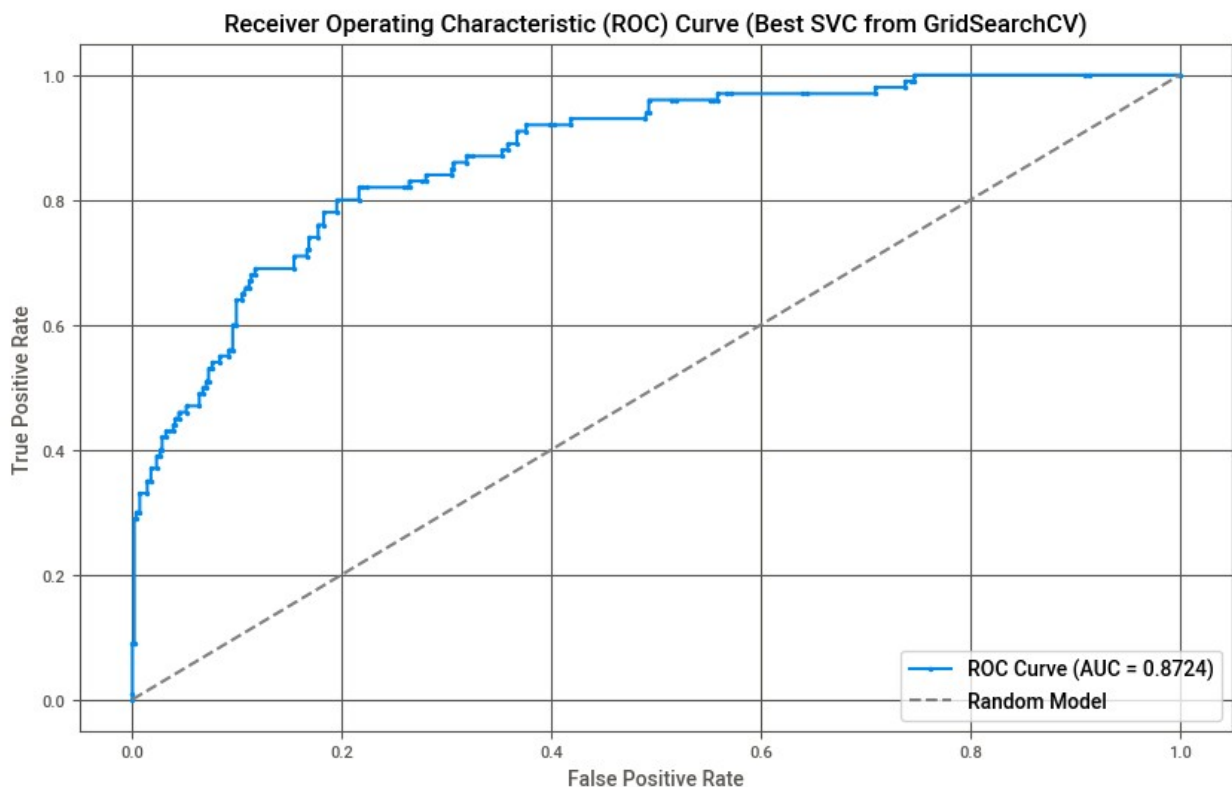
```

# Print AUC score
print(f"AUC Score (Best SVC): {auc_score:.4f}")

# Plot the ROC curve
plt.figure(figsize=(10, 6))
plt.plot(fpr, tpr, label=f"ROC Curve (AUC = {auc_score:.4f})",
marker='.')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray', label='Random Model')
plt.title("Receiver Operating Characteristic (ROC) Curve (Best SVC from GridSearchCV)")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.grid()
plt.show()

```

AUC Score (Best SVC): 0.8724



```

import numpy as np
from sklearn.metrics import roc_curve, roc_auc_score
import matplotlib.pyplot as plt

# Compute ROC curve and AUC score
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba_svcg)

```

```

auc_score = roc_auc_score(y_test, y_pred_proba_svcg)

# Compute False Negative Rate (FNR)
fnr = 1 - tpr

# Convert FPR and FNR to percentages
fpr_percentage = fpr * 100
fnr_percentage = fnr * 100

# Print AUC score
print(f"AUC Score (Best SVC): {auc_score:.4f}")

# Plot the DET curve
plt.figure(figsize=(10, 6))
plt.plot(fpr_percentage, fnr_percentage, label="DET Curve",
marker='.')

# Set logarithmic scale for both axes
plt.yscale("log")
plt.xscale("log")


# Add title and labels with percentages
plt.title("Detection Error Tradeoff (DET) Curve (Best SVC from
GridSearchCV)")
plt.xlabel("False Positive Rate (%) [Log Scale]")
plt.ylabel("False Negative Rate (%) [Log Scale]")

# Add legend and grid
plt.legend(loc="upper right")
plt.grid(which='both', linestyle='--', linewidth=0.5)

# Show the plot
plt.show()

AUC Score (Best SVC): 0.8724

```



## Feature Importance using SHAPly values

```
import shap
import matplotlib.pyplot as plt

# Train the best SVC model from GridSearchCV
best_svc = gridsvc.best_estimator_

# Explain the model predictions using SHAP
explainer = shap.Explainer(best_svc.predict, X_train) # Use the
prediction function
shap_values = explainer(X_train)

# Feature importance plot
plt.title("Feature Importance (SHAP Values)")
shap.summary_plot(shap_values, X_train, plot_type="bar")

PermutationExplainer explainer: 1550it [1:17:29, 3.01s/it]
```

